




Introduction to Operating System

Process Synchronization

Chap6+7 Synchronization Tool & Example



Silberschatz, Galvin and Gagne ©2018

葉奕成 I-Cheng (Garrett) Yeh
2021/5, V9 for 10th ed.



content

- Overview
 - 1. Introduction
 - 2. System Structures
- Process Management
 - 3. Process Concept
 - 4. Multithreaded Programming
 - 5. Process Scheduling
 - 6. Synchronization Tools
 - 7. Synchronization Example

BACKGROUND

Review : Producer-Consumer Problem

- *producer* process produces information
- *consumer* process consume information

Producer-Consumer Problem – A2

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
int count = 0; //表示目前在buffer內的item數
```

Producer-Consumer Problem – A2

Producer code

```
item nextProduced;
while (true) {
    // Produce an item in nextProduced
    while (count == BUFFER SIZE)
        ; // do nothing -- no free buffers

    buffer[in] = nextProducer;
    in = (in + 1) % BUFFER SIZE;
    count++;
}
```

Consumer code

```
item nextConsumed;
while (true) {
    while (count == 0)
        ; // do nothing -- nothing to consume

    // remove an item from the buffer
    nextConsumed = buffer[out];
    out = (out + 1) % BUFFER SIZE;
    count--;
    return item;
}
```

Race Condition

- `count++` could be implemented as
 `register1 = count`
 `register1 = register1 + 1`
 `count = register1`
- `count--` could be implemented as
 `register2 = count`
 `register2 = register2 - 1`
 `count = register2`

Race Condition

- Initially, let counter = 5.

S0: P: **register1 = count** {register1 = 5}

S1: P: **register1 = register1 + 1** {register1 = 6}

S2: P: **count = register1** {count = 6}

S3: C: **register2 = count** {register2 = 6}

S4: C: **register2 = register2 - 1** {register2 = 5}

S5: C: **count = register2** {**count = 5**} – CORRECT!

Race Condition

- Initially, let counter = 5.

S0: P: **register1 = count** {register1 = 5}

S1: P: **register1 = register1 + 1** {register1 = 6}

S2: C: **register2 = count** {register2 = 5}

S3: C: **register2 = register2 - 1** {register2 = 4}

S4: P: **count = register1** {count = 6}

S5: C: **count = register2** {**count = 4**} - ???????

Race Condition

- Initially, let counter = 5.

S0: P: **register1 = count** {register1 = 5}

S1: P: **register1 = register1 + 1** {register1 = 6}

S2: C: **register2 = count** {register2 = 5}

S3: C: **register2 = register2 - 1** {register2 = 4}

S4: C: **count = register2** {count = 4}

S5: P: **count = register1** {**count = 6**} - ???????

Race Condition

- A Race Condition
 - Several processes concurrently access and manipulate the same data
 - A situation where the outcome of the execution depends on the particular order of process scheduling.
若沒提供互斥存取控制，可能會因為
process執行的順序不同而結果有所不同
 - We need to ensure only one process at a time can be manipulating the shared data
 - Process synchronization and coordination

THE CRITICAL-SECTION PROBLEM

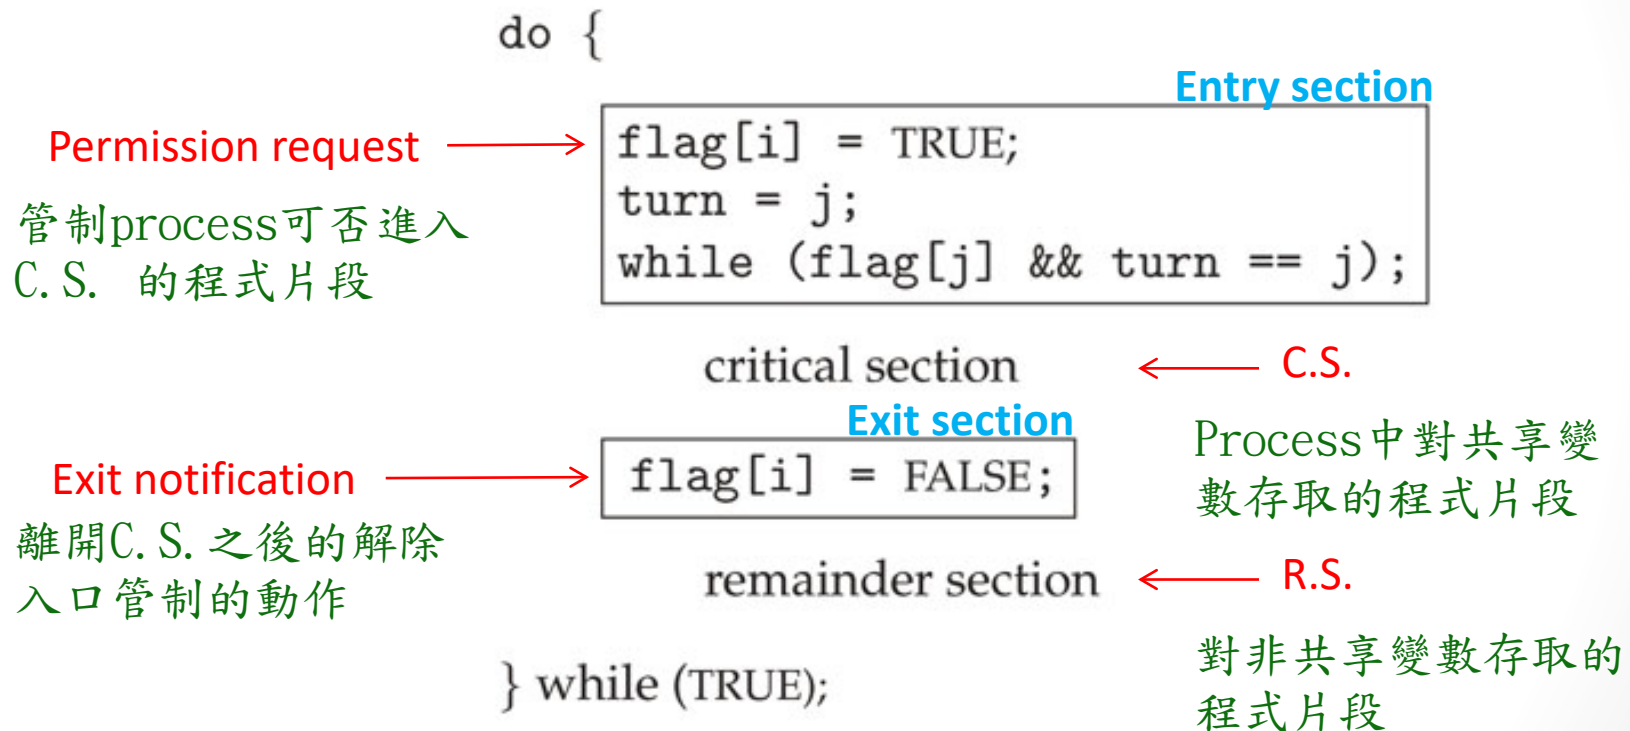
Framework for analysis of solutions

- Each process executes at nonzero speed but **no assumption** on the relative speed of N processes
 - = No assumption about order of interleaved execution
- Many CPUs may be present but memory hardware prevents simultaneous access to the same memory location
- For solutions: Processes may share some common variables to synchronize their actions.

The Critical-Section Problem

- Design a protocol that processes can use to cooperate.
 - Each process has a segment of code, called a **critical section**, whose execution must be **mutually exclusive**.
 - The section of code implementing this request is called the **entry section**.
 - The critical section (CS) might be followed by an **exit section**.
 - The remaining code is the **remainder section (RS)**.
 - mutually exclusive
 - one process is executing in its critical section, no other process is to be allowed to execute in its critical section.

The Critical-Section Problem



Requirements for a Valid Solution to Critical Section Problem

- Three Requirements

- **Mutual Exclusion**

- Only one process can be in its critical section.

- **Progress :** *程式要持續有進展

- If no process is executing in its CS and there exist some processes that wish to enter their CS...
 - (i) Only processes not in their remainder section can decide which will enter its critical section.

Process在非RS的位置可以決定誰可以進入CS

=> Process在RS的位置不能決定誰可以進入CS

=> 不想進入CS的Process不能阻礙其他Process進入他們各自的CS

- (ii) The selection cannot be postponed indefinitely.

- **No deadlock**

Requirements for a Valid Solution to Critical Section Problem

- Three Requirements
 - **Bounded Waiting**
 - After a process has made a request to enter its CS, there is a bound on the number of times that the other processes are allowed to enter their CS and before that request is granted.
 - A waiting process only waits for a bounded number of processes to enter their critical sections.
 - **No starvation**

Types of solutions

- Software solutions
 - Algorithms whose correctness does not rely on any other assumptions (see the framework)
- Hardware solutions
 - Rely on some special machine instructions
- Operation System solutions
 - Provide some functions and data structures to the programmer

Software solutions

- We consider first the case of 2 processes
 - Algorithm 1 and 2 are incorrect
 - Algorithm 3 is correct (**Peterson's algorithm**)
- Then we generalize to n processes
 - **Bakery algorithm**
- Notation
 - We start with 2 processes: P0 and P1
 - When presenting process P_i , P_j always denotes the other process ($i \neq j$)

Algorithm 1

- Shared Variable: **int turn**
 - initially **turn** = i or j;
 - *// means P_i or j can enter its critical section*
- Assumption
 - Every basic machine-language instruction is atomic.
假設基本的組合語言指令都是不可分割的
- Idea
 - Remember which process is allowed to enter its critical section, That is, *process i* can enter its critical section **if $\text{turn} = i$.**
*利用turn記錄誰可以進入C.S.
=> 誰有鑰匙(turn)誰就可以進入*

Algorithm 1

P_i

```
do
{
while(turn !=i) {no-op; //wait till turn==i
}
    C.S.
    turn=j; //release turn to j
    R.S.
}while(1)
```

Entry: **turn**不在手上，只好等待

Exit: 事情做完了
把**turn**交給另一個process

Algorithm 1

P₀

```
do
{
while(turn !=0) {no-op;
}
```

C.S.

```
turn=1;
```

R.S.

```
}while(1)
```

P₁

```
do
{
while(turn !=1) {no-op;
}
```

C.S.

```
turn=0;
```

R.S.

```
}while(1)
```

- When P₀ has a large RS and P₁ has a small RS....
- Problem: If turn=0, P₀ enter its CS and then its long RS (turn=1). P₁ enter its CS and then its RS (turn=0) and **tries again** to enter its CS: **request refused!**
- *P₁ has to wait that P₀ leaves its RS (not CS).

Algorithm 1 - step by step

- if $\text{turn} == 0$ and P1 is ready to enter its CS, P1 cannot do so, even though P0 may be in its RS. \rightarrow NO progress
 1. P0 enters, then exits from CS ($\text{turn} \Rightarrow 1$). Now, It is in its RS.
 2. Because $\text{turn} == 1$, P1 enters, then exits from CS ($\text{turn} \Rightarrow 0$).
P1 finishes this RS.
 3. P1 wants to enter CS **again**. Now $\text{turn} == 0$, P1 have to wait for P0 (even though P0 is in its RS).

Algorithm 1

- Mutual Exclusion

- Only one valid value for **turn**
->>> only one proc. has permit.
- satisfied

- Progress

- When $\text{turn} = i$, P_j is blocked even if P_i is running in the R.S.
*若 $\text{turn} = i$ ，但 P_i 不想進入，則 P_j 也無法進入

- No satisfies (i) in Progress requirement

- Conclusion: Satisfy mutual exclusion, but not progress

WRONG

P_i

```
do
{
while (turn  $\neq i$ ) {no-op;
}

C.S.
turn =  $j$ ;
R.S.
}while (1)
```


Algorithm 2

- Shared Variable
 - boolean **flag[2]**;
 - initially **flag[0] = flag[1] = false**;
 - When **flag[i] = true** >> P_i ready to enter its critical section
- Idea
 - Remember the state of each process.

利用flag記錄誰有意願進入C.S.

Algorithm 2

P_i

```
do {  
    flag[i] := true;  
    while (flag[j]){no-op;  
}  
    C.S.  
    flag[i] = false;  
    R.S.  
} while (1);
```

Entry:

先表達自己有意願

若對方有意願則禮讓對方做完

Exit:

事情做完表達自己沒意願

P_j

```
do {  
    flag[j] := true;  
    while (flag[i]){no-op;  
}  
    C.S.  
    flag[j] = false;  
    R.S.  
} while (1);
```

Algorithm 2

- May loop infinitely
 - Both of two proc. want to get into CS
若雙方都有意願，則雙方都在禮讓對方
導致都無法繼續往下執行
- If changes the order to be
 - while (flag[j]) {no-op;} //wait first**
 - flag[i] := true; //set flag later..**
 - ... C . S .**
- then the condition of mutual exclusion will **not** hold.

Algorithm 3: Peterson's Solution

- Combined shared variables of algorithms 1 and 2.
- Shared Variable
 - 利用 `turn` 記錄誰可以進入 C. S.
 - 利用 `flag` 記錄誰有意願進入 C. S.
 - \Rightarrow 有意願且有鑰匙 (`turn`) 的 process 可進入 C. S.
- `int turn;`
 - `turn = i // Pi can enter its critical section`
- `boolean flag[2];`
 - `flag[0] = flag[1] = false`
- “**turn**” will be set for both i and j simultaneously, but only one (`turn = i`) or (`turn = j`) will last.
- Meets all three requirements; solves the critical-section problem for two processes.

Algorithm 3

P_i

```
do {
    flag[i] := true;
    turn = j;
    while (flag [j] and turn == j) {no-op;
    }

    C.S.
    flag[i] = false;
    R.S.
} while (1);
```

Entry :

先表達自己有意願

再禮讓對方先做

若對方有意願則禮讓對方做完

若對方無意願則繼續往下做

Exit :

事情做完表達自己沒意願

就算flag[i]和[j]同時都是true，turn也會決定唯一的人選來進入C.S.
turn = j 的這個指令總是有一個人會先執行/一個人會後執行

Algorithm 3

P_j

```
do {  
    flag[j] := true;  
    turn = i;  
    while (flag [i] and turn == i) {no-op;  
    }  
    C.S.  
    flag[j] = false;  
    R.S.  
} while (1);
```

j

Algorithm 3

- Mutual exclusion
 - At the moment that P_i enters its CS
 - $\text{flag}[i]=\text{true}$ either $\text{flag}[j]=\text{false}$, or $\text{flag}[j]=\text{true}$ and $\text{turn}=i$
 - When P_j waits to enter *有兩種可能
 - If $\text{flag}[j]=\text{false}$, ($\text{flag}[i]=\text{true}$ & $\text{turn}=i$) prevents P_j from entering its CS until P_i exits
 - If $\text{flag}[j]=\text{true}$ and $\text{turn}=i$, the same (only P_i can change turn from i to j)

Algorithm 3

- Mutual exclusion
 - ...
- Progress
 - suppose P_i wishes to enter CS
 - if $\text{flag}[j]=\text{false}$, P_i can enter
 - otherwise, *turn* allows either P_i or P_j to enter
- Bounded waiting: wait at most one entry of P_j

What about Process Failures?

- If all 3 criteria (ME, Progress, Bounded waiting) are satisfied, then a valid solution will provide robustness against failure of a process in its remainder section (RS)
 - since failure in RS is just like having an infinitely long RS
- However, **no valid solution** can provide **robustness** against a process failing in its critical section (CS)
 - A process P_i that fails in its CS does not signal that fact to other processes: for them P_i is still in its CS

Bakery Algorithm, Critical section for n processes

- Before entering its critical section, process receives a **number**. Holder of the smallest number enters the critical section.
- If processes P_i and P_j receive the same number, if $i < j$, then P_i is served first; else P_j is served first.
- The numbering scheme always generates numbers in increasing order of enumeration; i.e., 1,2,3,3,3,3,4,5...

*不會產生比場上數字還小的票卷

-- 後來的人號碼永遠比較大

-- 完成的人號碼歸零 (不採計大小比較)

Bakery Algorithm (cont.)

- Notation \leq lexicographical order (ticket #, process id #)
 - $(a,b) < (c,d)$: if $a < c$ or if $a == c$ and $b < d$
 - $\max(a_0, \dots, a_{n-1})$ is a number, k ,
 - such that $k \geq a_i$ for $i = 0, \dots, n-1$
- Shared data
 - boolean choosing[n];** //a state indicator
 - int number[n];** //a ticket

Data structures are initialized to **false** and **0** respectively.

Bakery Algorithm (cont.)

```
do { Take a ticket  
    choosing[i] = true;  
    number[i] = max(number[0], number[1], ..., number  
    [n - 1]) + 1;  
    choosing[i] = false;  
    for (j = 0; j < n; j++) {  
        while (choosing[j]) ; /* Wait for the choosing of Pj  
        while ((number[j] != 0) && (number[j], j) <  
            number[i], i)) ; // smallest first  
    }  
    critical section  
    number[i] = 0;  
    remainder section  
} while (1);
```

Bakery Algorithm

- Why $\text{number}[i] = \text{number}[j]$?
 - P_i : $\text{choosing}[i] = \text{true}$;
 - P_j : $\text{choosing}[j] = \text{true}$; The largest ticket is **k** for now

- P_i : $\text{Rigister1} = \max(\text{number}[0] \dots \text{number}[n-1]) + 1$
- P_j : $\text{Rigister2} = \max(\text{number}[0] \dots \text{number}[n-1]) + 1$
- P_i : $\text{number}[i] = \text{Rigister1} // == k+1$
- P_j : $\text{number}[j] = \text{Rigister2} // == k+1$

Bakery Algorithm

- When we remove `while (choosing[j]){no-op} ...`
 - No satisfies mutual exclusion

default:

`number[0..n-1]=0,`

`Pid i < Pid j,`

`Pi, Pj` prepare to enter the C.S.

`Pi : choosing[i] = true;`

`Pj : choosing[j] = true;`

The largest ticket is **k** for now

`Pi : number[i] = max(number[0] ... number[n-1])+1 //not assign yet!!`

`Pj : number[j] = max(number[0] ... number[n-1])+1 // ==k+1`

`Pj : choosing[j] = false`

`Pj : Pj in C.S`

`Pi : number[i] = k+1 // complete the assignment step now...`

`Pi : Pi in C.S !!!`

=> No satisfies mutual exclusion

Bakery Algorithm (cont.)

- Key for showing the correctness
 - if P_i in CS, all other P_k has
 - number[k] = 0, or
 - (number[i], i) < (number[k], k)
- mutual exclusion: OK
- progress: OK (smallest first)
- bounded waiting: OK
 - Note that processes enter their CSs in a FCFS basis.
 - How many times ? n

Drawbacks of Software Solutions

- Processes that are requesting to enter in their critical section are busy waiting (consuming processor time needlessly)
- If Critical Sections are long, it would be more efficient to block processes that are waiting...

Synchronization Hardware

HARDWARE SUPPORT FOR SYNCHRONIZATION

Hardware Solutions ?

- Solution to Critical-section Problem Using Locks
- A process ...
 - **Acquire** a lock before entering a critical section
 - **Releases** the lock when it exits the critical section

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (TRUE);
```

Hardware Solutions: Interrupt Disabling

- On a uniprocessor:
 - Mutual exclusion is preserved but **efficiency of execution is degraded**: while in CS, we cannot interleave execution with other processes that are in RS.
- On a multiprocessor: mutual exclusion is not preserved
 - CS is now atomic but not mutually exclusive.
 - Potential impacts on interrupt-driven system clocks.

Process P_i :

repeat

disable interrupts

critical section

enable interrupts

remainder section

forever

Hardware Solutions:

Special Machine Instructions

- Normally, access to a memory location excludes other access to that same location
 - Extension: designers have proposed machine instructions that perform 2 actions atomically (indivisible) on the same memory location (ex: reading and writing)
 - The execution of such an instruction is also mutually exclusive (even with multiple CPUs)
- They can be used to provide mutual exclusion but need to be complemented by other mechanisms to satisfy the other 2 requirements of the CS problem (We still need to avoid starvation and deadlock)

Synchronization Hardware (1)

- Having the support of some simple hardware instructions, the CS problem can be solved very easily and efficiently.
- The CS problem occurs because the modification of a shared variable of a process may be interrupted.
- Two common hardware instructions that execute

atomically

- *Test-and-Set*
- *Swap*

*Atomic = non-interruptable

Atomically executed:

執行過程中不會被中斷

可在一個memory cycle內完成工作

Synchronization Hardware (2)

- Test and modify the content of a word atomically

```
boolean TestAndSet(boolean &target) {
```

```
    boolean rv = target; //backup original value
```

```
    target = true; //set target = true
```

```
    return rv; //send out original value
```

```
}
```


No
Preemption



Mutual Exclusion with Test-and-Set

- Shared data:
boolean lock = false;

- Process P_i
do {
 while (TestAndSet(lock)) ; no-op
 critical section
 lock = false;
 remainder section
}



For example:

Time 0: P_i in: TestAndSet(&lock) //return false and set lock =true > P_i in C.S

Time 1: P_j in: TestAndSet(&lock) //return true and set lock =true > P_j in loop...

Mutual Exclusion with Test-and-Set

```
do {  
    while (TestAndSet(lock)) ;  
        critical section  
    lock = false;  
        remainder section  
}
```

- Mutual Exclusion

- 有鑰匙才能進 -> ok
- 執行過程中不會被中斷 :
 - *可在一個memory cycle內完成工作

- Progress

- 不想進入C. S就不拿鑰匙 -> ok
- 有限時間內可決定鑰匙給誰 -> ok

- Bounded Waiting

- P_i 離開後又想進入C. S, P_i 直接執行TestAndSet (比 P_j 早/ P_j 不在等待時) 則 P_i 又進入C. S.
- Starvation ... ??? (有可能一直把持著鑰匙不放)

Extend Test-and-Set to N Processes

false: 無意願進入C.S 或即將進入C.S
true: 有意願 & 正在等待進入C.S

- Shared data:

```
boolean lock = false;  
boolean waiting[0..n-1] //default : false  waiting
```

- Process P_i

```
do {
```

```
    waiting[i]=true; //表達意願
```

```
    key = true;
```

```
    while(waiting[i] and key) {key = TestAndSet(&lock);}
```

```
    waiting[i]=false;
```

```
        critical section
```

找下一個
有意願的

```
    { j=(i+1) mod n;
```

```
      while (j != i and (not waiting[j])) {j=(j+1) mod n;}
```

```
      if (j == i) {lock = false;} // j=i 表所有人無意願
```

```
      else {waiting[j]=false;} //讓下一個有意願的離開while
```

```
        remainder section
```

```
    }
```

key : boolean
PID j : 0 ... n-1

Mutual Exclusion with Test-and-Set

- Mutual Exclusion:

靠別人幫忙

- P_i can enter CS only if either $waiting[i] == false$ or $key == false$ P_i 必須自己先搶到
- key becomes false only if TestAndSet is executed
 - First process to execute TestAndSet find $key == false$; others wait ...
- $waiting[i]$ becomes false **only if other process leaves CS**
 - Only one $waiting[i]$ is set to false (the next available one)

- Progress

- 不想進入C.S就不拿鑰匙 $\rightarrow waiting[i] = false \rightarrow ok$
- 有限時間內可決定鑰匙給誰 $\rightarrow ok$

Mutual Exclusion with Test-and-Set


- Bounded Waiting
 - 若所有 process 皆想進入 C.S
T0: P0 : waiting[1]=false;
T1: P1 : in C.S....
P0 想再進入 C.S, 需等待 Pn-1 完成後才能進入 C.S
 - No Starvation : 一定會換下一個人

Synchronization Hardware

TYPE 2

- Atomically swap two variables.

```
void Swap(boolean &a, boolean &b) {  
    boolean temp = a;  
    a = b;  
    b = temp;  
}
```



No Preemption
Swap(lock, key);

Mutual Exclusion with Swap

TYPE 2

- Shared data (initialized to **false**):
 boolean lock;
 boolean waiting[n];
- Process P_i
 do {
 key = true;
 while (key == true)
 Swap(lock, key);
 critical section
 lock = false;
 remainder section
 }

Mutual Exclusion with Swap

```
do {  
    key = true;  
    while (key == true)  
        Swap(lock, key);  
    critical section  
    lock = false;  
    remainder section  
}
```

- Mutual Exclusion

- 有鑰匙(& lock==false)才能進 -> ok
- 交換鑰匙過程中不會被中斷：
*可在一個memory cycle內完成工作


- Progress

- 不想進入C. S就不拿鑰匙 -> ok
- 有限時間內可決定鑰匙給誰 -> ok

- Bounded Waiting

- P_i 離開後又想進入C. S > P_i 比 P_j 早下Swap, 則 P_i 又進入C. S.
- Starvation ... ??? (把持著鑰匙不放)

Mutual Exclusion with Swap

- Shared data:
boolean lock = false;  false:無意願進入C.S 或即將進入C.S
boolean waiting[0..n-1] //default : false true:有意願 & 正在等待進入C.S
waiting
- Process P_i
do {

waiting[i]=true; //表達意願

key = true;
while(waiting[i] and key) {Swap(lock,key);}
waiting[i]=false;
critical section
找下一個有意願的 { j=(i+1) mod n;
while (j != i and (not waiting[j])) {j=(j+1) mod n;}
if (j ==i) {lock = false;} // j=i 表所有人無意願
else {waiting[j]=false;} //讓下一個有意願的離開while
remainder section
}
}

key : boolean
PID j : 0 ... n-1

SEMAPHORE

Overview -- Semaphore

- A high-level solution for more complex problems.
 - Synchronization tool (provided by the OS) that does not require busy waiting
- A semaphore S is an integer variable that, apart from initialization, **can only be accessed** through **2 atomic and mutually exclusive** operations:
 - **wait(S)** – P operation *proberen* (荷蘭文:測試)
 - **signal(S)** – V operation *verhogen* (荷蘭文:增加)
- To avoid busy waiting: when a process has to wait, it will be put in a blocked queue of processes waiting for the same event

Semaphore_{cont.}

- **wait(S)**: **decrements** the **semaphore value**. If the value becomes **negative**, then the process (who is executing the wait(s)) is **blocked**
- **signal(S)**: **increments** the semaphore value. If the value is not positive, then a process blocked by a wait(s) operation is **unblocked**.
- wait and signal are assumed to be **atomic**
 - They cannot be interrupted and each routine can be treated as an indivisible step

不能分割的

Semaphore

```
wait(S) {  
    while (S <= 0) ; // no-op  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

Semaphores – Usages

- Handling “Critical Sections Problem”
 - default : Semaphore S = 1;
- Enforce the requirement of precedence
 - default : Semaphore S = 0;

Semaphores – Usages

- Critical Sections
 - Shared Variable
 - semaphore mutex= 1; // allow first process go into C.S.

P_i

```
do
{
    wait(mutex);
    C.S
    signal(mutex);
    R.S
}while(TRUE)
```

Semaphores – Usages

- Precedence Enforcement
 - Shared Variable
 - semaphore synch= 0;
 - Procedure
 - S1 (in P1) -> S2 (in P2)

P₁

```
S1;  
  signal(synch);
```

P₂

```
  wait(synch);  
  
S2;
```

Semaphores – Usages

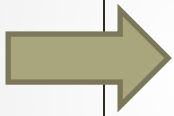
- Counting semaphore = General semaphore
 - This initialized number represents a *number of available resources*.
 - (Non-negative) integer value can range over an unrestricted domain.
- Binary semaphore
 - Integer value can range only between 0 and 1
 - can be simpler to implement
 - Also known as **mutex locks**
- We can implement a counting semaphore S as a binary semaphore.

Semaphores – Usages

- Using Binary semaphore implement a counting semaphore

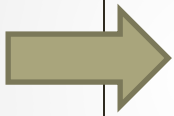
```
semaphore s1= 1 ; // protect c
semaphore s2= 0 ; // wait for ( c <=0 )
int c = 1 ; //counter
```


Semaphores – Usages



```
WAIT(S) // check&use resource
{
    wait(s1); // protect C
    c--;
    if (c < 0) {
        signal(s1);
        wait(s2); // block ...
    }
    else
        signal(s1);
}
```

Semaphores – Usages



```
SIGNAL(S) //release resource
{
    wait(s1);
    c++;
    if (c <= 0)
        signal(s2); //wakeup others
    signal(s1);
}
```

A red bracket is drawn on the right side of the code block, spanning from the 'wait(s1);' line down to the 'signal(s1);' line, indicating a critical section.

Semaphores – Usages

- P1 in
 - P1 get into C.S

```
semaphore s1= 1 ; //1->0->1
semaphore s2= 0 ;
int c = 0 ; // 1->0
```

```
WAIT (S)
{
    wait(s1);
    c--;
    if (c < 0) {
        signal(s1);
        wait(s2);
    } else signal(s1);
}
```

1->0

Protect C

0->1

```
do
{
    WAIT(S);
    C.S
    SIGNAL(S);
    R.S
}while(TRUE)
```

```
SIGNAL (S)
{
    wait(s1);
    c++;
    if (c <= 0)
        signal (s2); /* wakeup */
    signal (s1);
}
```

Semaphores – Usages

- P2 in
 - P2 wait

```
semaphore s1= 1 ;  
semaphore s2= 0 ;  
int c = -1 ;
```

```
WAIT (S)  
{  
    wait(s1); /* protect C */  
    c--;  
    if (c < 0) {  
        signal(s1);  
        wait(s2);  
    } else signal(s1);  
}
```

← s2==0

```
do  
{  
    WAIT(S);  
    C.S  
    SIGNAL(S);  
    R.S  
}while(TRUE)
```

```
SIGNAL (S)  
{  
    wait(s1);  
    c++;  
    if (c <= 0)  
        signal (s2); /* wakeup */  
    signal (s1);  
}
```

Semaphores – Usages

- **P3 in**
 - **P3 wait (same as P2)**

|c| : mean # of process wait

```
semaphore s1= 1 ;
semaphore s2= 0 ;
int c = -2 ;
```

```
WAIT (S)
{
    wait(s1); /* protect C */
    c--;
    if (c < 0) {
        signal(s1);
        wait(s2);
    } else signal(s1);
}
```

← s2==0

```
do
{
    WAIT(S);
    C.S
    SIGNAL(S);
    R.S
}while(TRUE)
```

```
SIGNAL (S)
{
    wait(s1);
    c++;
    if (c <= 0)
        signal (s2); /* wakeup */
    signal (s1);
}
```


Semaphores – Usages

- **P1 out**



```
semaphore s1= 1 ;  
semaphore s2= 1 ; // s2++ : 0->1  
int c = -1 ; // c++
```

```
WAIT (S)  
{  
    wait(s1); /* protect C */  
    c--;  
    if (c < 0) {  
        signal(s1);  
        wait(s2);  
    } else signal(s1);  
}
```

```
do  
{  
    WAIT(S);  
    C.S  
    SIGNAL(S);  
    R.S  
}while(TRUE)
```



```
SIGNAL (S)  
{  
    wait(s1);  
    c++;  
    if (c <= 0)  
        signal (s2); /* wakeup */  
    signal (s1);  
}
```



Semaphores – Usages

- P2 in C.S
 - Only P3 is waiting

```
semaphore s1= 1 ;  
semaphore s2= 0 ; // s2-- : 0->1  
int c = -1 ;
```

```
WAIT (S)  
{  
    wait(s1); /* protect C */  
    c--;  
    if (c < 0) {  
        signal(s1);  
        wait(s2);  
    } else signal(s1);  
}
```

← s2--

```
do  
{  
    WAIT(S);  
    C.S  
    SIGNAL(S);  
    R.S  
}while(TRUE)
```

```
SIGNAL (S)  
{  
    wait(s1);  
    c++;  
    if (c <= 0)  
        signal (s2); /* wakeup */  
    signal (s1);  
}
```

Semaphores – Usages

- P4 in

- P4(&P3) wait

|c| : mean # of process wait

```
semaphore s1= 1, //1->0->1  
semaphore s2= 0 ;  
int c = -2 ;
```

```
WAIT (S)  
{  
    wait(s1); /* protect C */  
    c--;  
    if (c < 0) {  
        signal(s1);  
        wait(s2);  
    } else signal(s1);  
}
```

New P comes here.

s2==0

```
do  
{  
    WAIT(S);  
    C.S  
    SIGNAL(S);  
    R.S  
}while(TRUE)
```

```
SIGNAL (S)  
{  
    wait(s1);  
    c++;  
    if (c <= 0)  
        signal (s2); /* wakeup */  
    signal (s1);  
}
```


Semaphores - Implementation

- The main **disadvantage** of **previous semaphore** and **other schemes** is that they all require ***busy waiting***.
 - It wastes CPU cycles that some other process might be able to use productively in multiprogramming system.

Semaphores - Implementation

- This type of semaphore is also called a **spinlock**
盤旋鎖
 - A **Busy-Waiting** Semaphore
 - while ($S \leq 0$) causes the wasting of CPU cycles
- Advantage : (in multiprocessor system)
 - When locks are held for a short time, spinlocks are useful since no context switching is involved.
 - *One thread can “spin” on one processor while another thread performs its critical section on another processor.)

因不用context switching，所以若只有很短的時間就可以離開while，spinlock是有用的

Semaphores - Implementation

- Hence, in fact, a semaphore is a record (structure)
- When a process must **wait** for a semaphore S, it can **block itself and put on the semaphore's queue**
 - The queue links PCB
- The signal operation removes (according to a fair policy like FIFO) one process from the queue and puts it in the list of ready processes
- Block and wakeup change process state – they are basic system calls
 - **Block:** from running to waiting
 - **Wakeup:** from waiting to ready

Semaphores - Implementation

- Semaphores with Block-Waiting
 - No busy waiting from the entry to the critical section
 - With each semaphore there is an associated **waiting queue**.
 - Each semaphore has two data items:
 - value (an integer)
 - pointer to a list of process

```
typedef struct {  
    int value;  
    struct process *list;  
} semaphore ;
```

Semaphores - Implementation

- Semaphores with Block-Waiting, two operations:
 - **block**
 - place the process who invoking the operation on the appropriate waiting queue.
 - **wakeup** 將process自己放到適當的waiting queue
 - remove one of processes in the waiting queue and place it in the ready queue. 將某process從waiting queue取出並放到ready queue

Semaphores - Implementation

- Semaphores with Block-Waiting

```
wait(semaphore *S) {  
    S->value--; |S->value| : mean # of process wait  
    if (S->value < 0) {  
        add this process to S->list;  
        block() ;  
    } // S->value >= 0  
}
```

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P) ;  
    }  
}
```

Semaphores - Observations

- When $S.value \geq 0$: the number of **processes that can execute** $wait(S)$ **without** being blocked = $S.value$ = # or resource
- When $S.value < 0$: the number of **processes waiting** on S is = $|S.value|$
- Atomicity and mutual exclusion: no 2 process can be in $wait(S)$ and $signal(S)$ (on the same S) at the same time (even with multiple CPUs)
- Hence the blocks of code defining $wait(S)$ and $signal(S)$ are, in fact, critical sections, too.

Semaphores - Observations_{cont.}

- Wait and Signal must be executed atomically (mutual execution)
- The critical sections defined by wait(S) and signal(S) are very short: **typically 10 instructions**
- Solutions:
 - uniprocessor: disable interrupts during these operations (ie: for a very short period). This does not work on a multiprocessor machine.
 - multiprocessor: use previous software or hardware schemes. The amount of busy waiting should be small.

Deadlocks and Starvation

- Deadlock – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.

Let S and Q be two semaphores initialized to 1

P0: wait(S);	P1: wait(Q);
wait(Q);	wait(S);
...	...
signal(S);	signal(Q);
signal(Q);	signal(S);

One more thing...

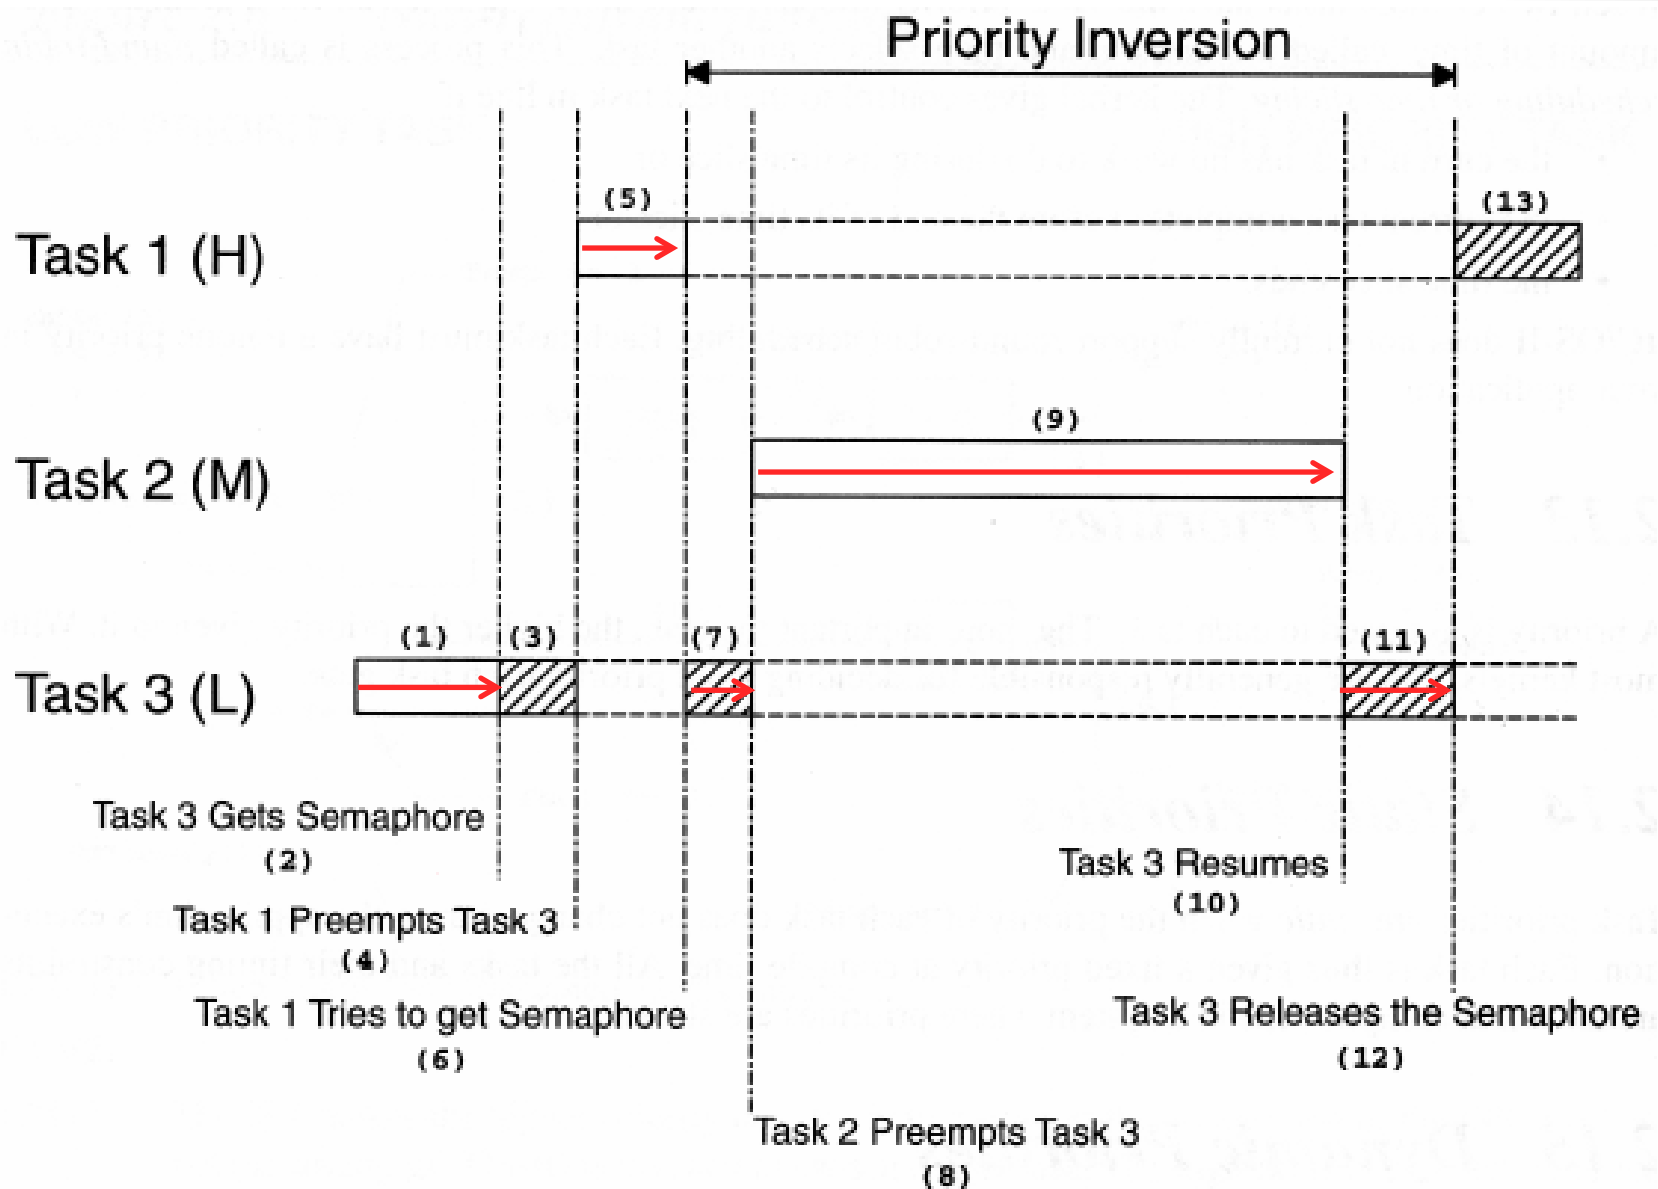
- The queuing strategy can be arbitrary, but there is a restriction for the **bounded waiting** requirement.
 - For example, **FIFO** is a common choice.
- Starvation (or indefinite blocking) may occur if we add and remove processes from the semaphore queue in **LIFO** order

ISSUE: Priority Inversion

- **Priority Inversion** 優先權倒置
 - A higher-priority task is blocked by a lower-priority task due to some resource access conflict.
 - Scheduling problem when lower-priority process holds a lock needed by higher-priority process
 - The situation will be worse if the lower – priority process is preempted by other higher-priority processes.

高優先權需要的資源被低優先權把持住，高優先權的process等待。若之後有一個次高優先權的process進來，並搶走低優先權process的CPU，則發生高優先權等待次高優先權process的情況

ISSUE: Priority Inversion



ISSUE: Priority Inversion

- **Priority-inheritance protocol** 優先權繼承協定
 - all processes that are accessing resources needed by a higher-priority process inherit the higher priority until they are finished with the resources

把持高優先權所需資源的低優先權process，將繼承高優先權process的優先權，直到釋放資源為止

Classical Problems of Synchronization

- There are three classical problem
 - Bounded-Buffer Problem (Producer-Consumer Problem)
 - Readers and Writers Problem
 - Dining-Philosophers Problem
- The three problems are important, because ...
 - They are examples for a large class of concurrency-control problems
 - They has been used for testing nearly every newly proposed synchronization scheme.

Classical Problems of Synchronization

How to solve these problems by using semaphore?

Producer-consumer (bounded-buffer) problem:


* A pool consists of n buffers (each holds a item)

semaphore full, empty, mutex;

mutex(1): control of mutual exclusion

empty(n), full(0): the number of empty and full buffers

Bounded-Buffer Problem



- Shared variable
 - N buffers, each can hold one item
 - Semaphore **mutex** = 1;
 - Semaphore **full** = 0;
 - 0 : no item
 - >0 : many items
 - Semaphore **empty** = n ;
 - 0 : buffer full
 - >0 : buffer is not full

Bounded-Buffer Problem

- Producer

等待buffer 有空間，若有空間，則產生item，空間數-1

```
do {  
    // produce an item in nextp  
    wait (empty); // if empty = 0 ; buffer full; wait  
    wait (mutex);  
    // add the item to the buffer  
    signal (mutex);  
    signal (full); // full +1 提醒consumer , buffer已經有item  
} while (TRUE);
```

Bounded-Buffer Problem

- Consumer

等待buffer 有item，若有item
，則消耗item，item數-1

```
do {  
    wait (full);  
    wait (mutex);  
    // remove an item from buffer  
    signal (mutex);  
    signal (empty); // empty +1 提醒producer, buffer已經有空間  
    // consume the item in nextc  
} while (TRUE);
```

Bounded-Buffer Problem

Producer

```
do {  
    // produce an item in nextp  
    wait (empty);  
    // if empty = 0 ; buffer full; wait  
    wait (mutex);  
    // add the item to the buffer  
    signal (mutex);  
    signal (full); // full +1  
} while (TRUE);
```

Consumer

```
do {  
    wait (full);  
    wait (mutex);  
    // remove an item from buffer to nextc  
    signal (mutex);  
    signal (empty); // empty +1  
    // consume the item in nextc  
} while (TRUE);
```

Readers-Writers Problem (1)

- each writer is required to have exclusive access to the shared object.
- At the same time, either several reads, or only a write

- Shared data

semaphore mutex, wrt;

readcount: (0, initially; # of current readers)

mutex: (1, initially; mutual exclusion for *readcount*)

wrt: (1, initially)

- mutual exclusion for writers.
- It also is used by the *first-in* and *last-out* readers

Readers-Writers Problem (2)

- If a writer is in the CS and n readers are waiting, then
 - one is queued on *wrt*
 - the other $n-1$ are queued on *mutex*
- When a writer executes *signal(wrt)*, either the waiting readers or a single writer are resumed. The selection is made by the scheduler.

Readers and Writers Problem (4)

for the first solution

writer

```
wait(wrt);  
...  
writing is performed  
...  
signal(wrt);
```

reader



first-in

```
wait(mutex);  
readcount++;  
if (readcount == 1) wait(wrt);  
signal(mutex);  
...  
reading is performed  
...  
wait(mutex);  
readcount--;  
if (readcount == 0) signal(wrt);  
signal(mutex);
```

update
readcount

Dining-Philosophers Problem

- Five philosophers, either thinking or eating
 - Each philosopher must pick up one chopstick beside him/her at a time
 - When two chopsticks are picked up, the philosopher can eat.
- Shared variable
 - Bowl of rice (data set)
 - Semaphore **chopstick** [5] = 1;



Dining-Philosophers Problem

philosopher i

```
do {  
    wait(chopstick[i]);  
    wait(chopstick[(i+1) % 5]);  
    ...  
    eat  
    ...  
    signal(chopstick[i]);  
    signal (chopstick[(i+1) % 5]);  
    ...  
    think  
    ...  
} while(1);
```

看左邊的筷子是否可拿

看右邊的筷子是否可拿

get chopsticks

left

right

free chopsticks

left

right

deadlock !

Dining-Philosophers Problem

- Deadlock occur
 - If all philosophers pick up their right one simultaneously.
- Common Solutions for Deadlocks:
 1. Allow at most **four** philosophers to be sitting simultaneously at the table.
 2. Pick up two chopsticks simultaneously.
 3. Order their behaviors,
 - e.g., odds pick up their right one first, and evens pick up their left one first.
- Besides deadlock, any satisfactory solution to the DPP problem must avoid the problem of *starvation*.

Dining-Philosophers Problem

- Solutions to Deadlocks
 - Semaphore **num = 4;**

```
do {  
    wait (num);  
    wait ( chopstick[i] );  
    wait ( chopStick[ (i + 1) % 5] );  
    // eat  
    signal ( chopstick[i] );  
    signal ( chopstick[ (i + 1) % 5] );  
    signal (num);  
    // think  
} while (TRUE);
```

MONITOR

Monitor

- Semaphores provide a convenient and effective mechanism for process synchronization.
- However, using them **incorrectly** can result in timing errors that are difficult to detect.
 - Since these errors happen **only if particular execution sequences** take place and these sequences do not always occur.
 - For example, exchange `wait()` and `signal()` in the client code, or simply omit them...

Monitor


An **abstract data type**—or **ADT**—encapsulates data with a set of functions to operate on that data that are independent of any specific implementation of the ADT.

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization—the ***monitor type***.
 - Only one process may be active within the monitor at a time
- A ***monitor type*** is an ADT that includes a set of operations that are provided with mutual exclusion within the monitor.
 - The monitor type also declares the variables whose values define the state of an instance of that type, along with the bodies of functions that operate on those variables.

Monitor: A High-level Language Constructs

- The representation of a monitor type consists of
 - **declarations of variables** whose values define the state of an instance of the type
 - **procedures or functions** that implement operations on the type.
- A procedure within a monitor can access only variables defined in the monitor.
 - The local variables of a monitor can be used only by the local procedures.
- The monitor construct ensures that only one process at a time can be active within the monitor.
 - Consequently, the programmer does not need to code this synchronization constraint explicitly

Monitor



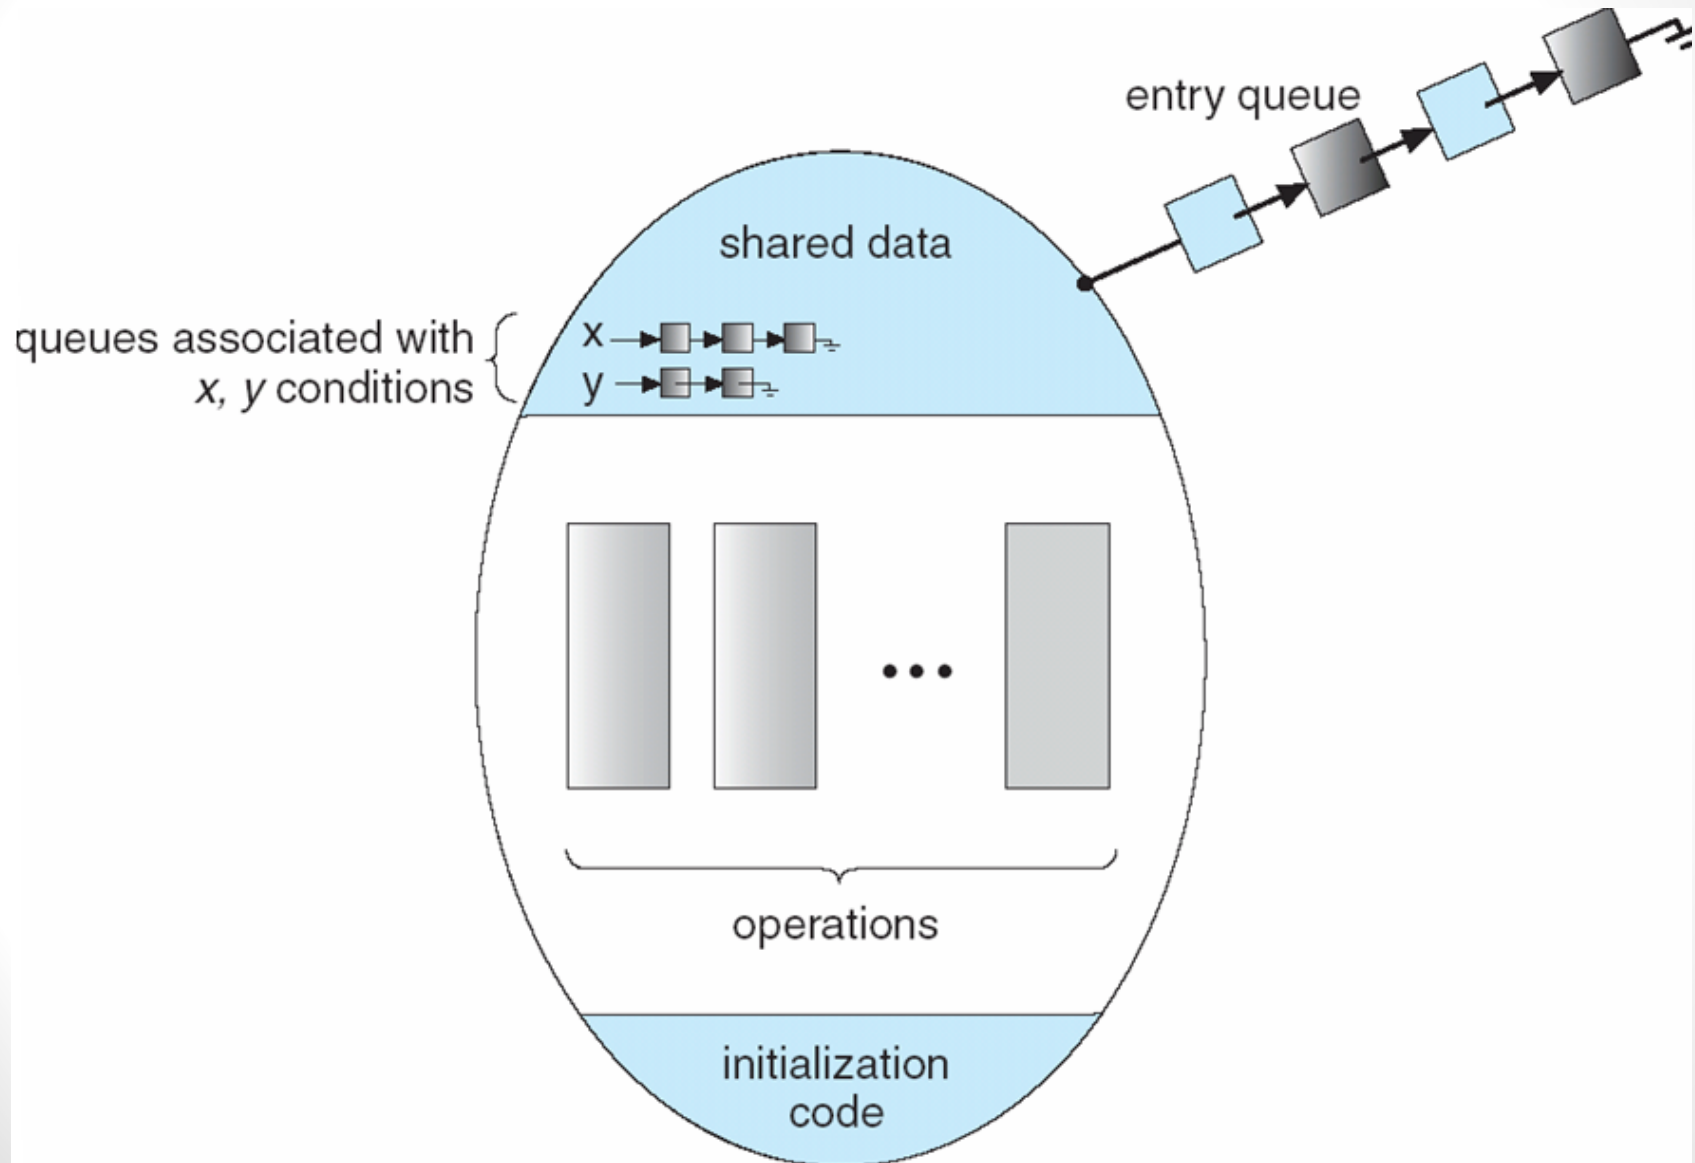
```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }
    ...
    procedure Pn (...) {.....}

    Initialization code ( ....) { ... }
    ...
}
}
```

Monitor

- We define additional synchronization mechanisms which provided by the **condition construct**.
 - A programmer who needs to write a tailor-made synchronization scheme can define one or more variables of type *condition*:
 - condition x, y;
- Condition variables
 - x.wait () : a process that invokes the operation is suspended.
 - x.signal () : Resumes exactly one of processes (if any) that invoked x.wait ()
 - Contrast this operation with the signal() operation associated with semaphores, which always affects the state of the semaphore.

Monitor with Condition Variables



Condition Variables Choices


- If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`, what should happen next?
 - Both Q and P cannot execute in parallel. If Q is resumed, then P must wait
- Options include
 - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
 - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
 - Both have pros and cons – language implementer can decide
 - Monitors implemented in Concurrent Pascal compromise
 - P executing signal immediately leaves the monitor, Q is resumed
 - Implemented in other languages including C#, Java

Resuming Processes within a Monitor

- If several processes queued on condition variable x , and $x.\text{signal}()$ is executed, which process should be resumed?
- FCFS frequently not adequate
- **conditional-wait** construct of the form $x.\text{wait}(c)$
 - Where c is **priority number**
 - Process with lowest number (highest priority) is scheduled next

Example with priority numbers

- Allocate a single resource among competing processes using **priority numbers** that **specify the maximum time a process plans to use** the resource




```
R.acquire(t) ;  
    ...  
    access the resource ;  
    ...  
R.release ;
```

- Where R is an instance of type **ResourceAllocator**


A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = true;
    }
    void release() {
        busy = FALSE;
        x.signal();
    }
    initialization code() {
        busy = false;
    }
}
```

the maximum time a process
plans to use the resource



The monitor allocates the resource to the
process that has the shortest time-allocation
request



Solution to Dining Philosophers

- We illustrate **monitor** concepts by presenting a deadlock-free solution to the dining-philosophers problem.
- This solution imposes the restriction that a philosopher may pick up her chopsticks only if both of them are available.

Solution to Dining Philosophers

- The solution
 - declare a monitor dp of type dining-philosophers

dp: dining-philosophers

- To eat, philosopher P_i performs

dp.pickup(i);

...

eat

...

dp.putdown(i);


Solution to Dining Philosophers

- Each philosopher i invokes the operations `pickup()` and `putdown()` in the following sequence

```
DP DiningPhilosophers;  
DiningPhilosophers.pickup(i);  
    ... eat ...  
DiningPhilosophers.putdown(i);  
    ... think ...
```

- Philosopher i can set the variable `state[i] = EATING` only if her two neighbors are not eating.

Solution to Dining Philosophers



monitor DP

{


```
enum { THINKING; HUNGRY, EATING) state [5] ;  
condition self [5];
```

```
void pickup (int i) { /*肚子餓撿起筷子*/  
    state[i] = HUNGRY;  
    test(i); /*測試是否可以拿起筷子*/  
    if (state[i] != EATING) self [i].wait;  
}
```

```
void putdown (int i) { /*吃飽放下筷子*/  
    state[i] = THINKING;  
    // test left and right neighbors  
    test((i + 4) % 5); /*詢問旁邊是否要吃飯 (把他們叫起來測試)*/  
    test((i + 1) % 5); /*詢問旁邊是否要吃飯*/  
}
```

....

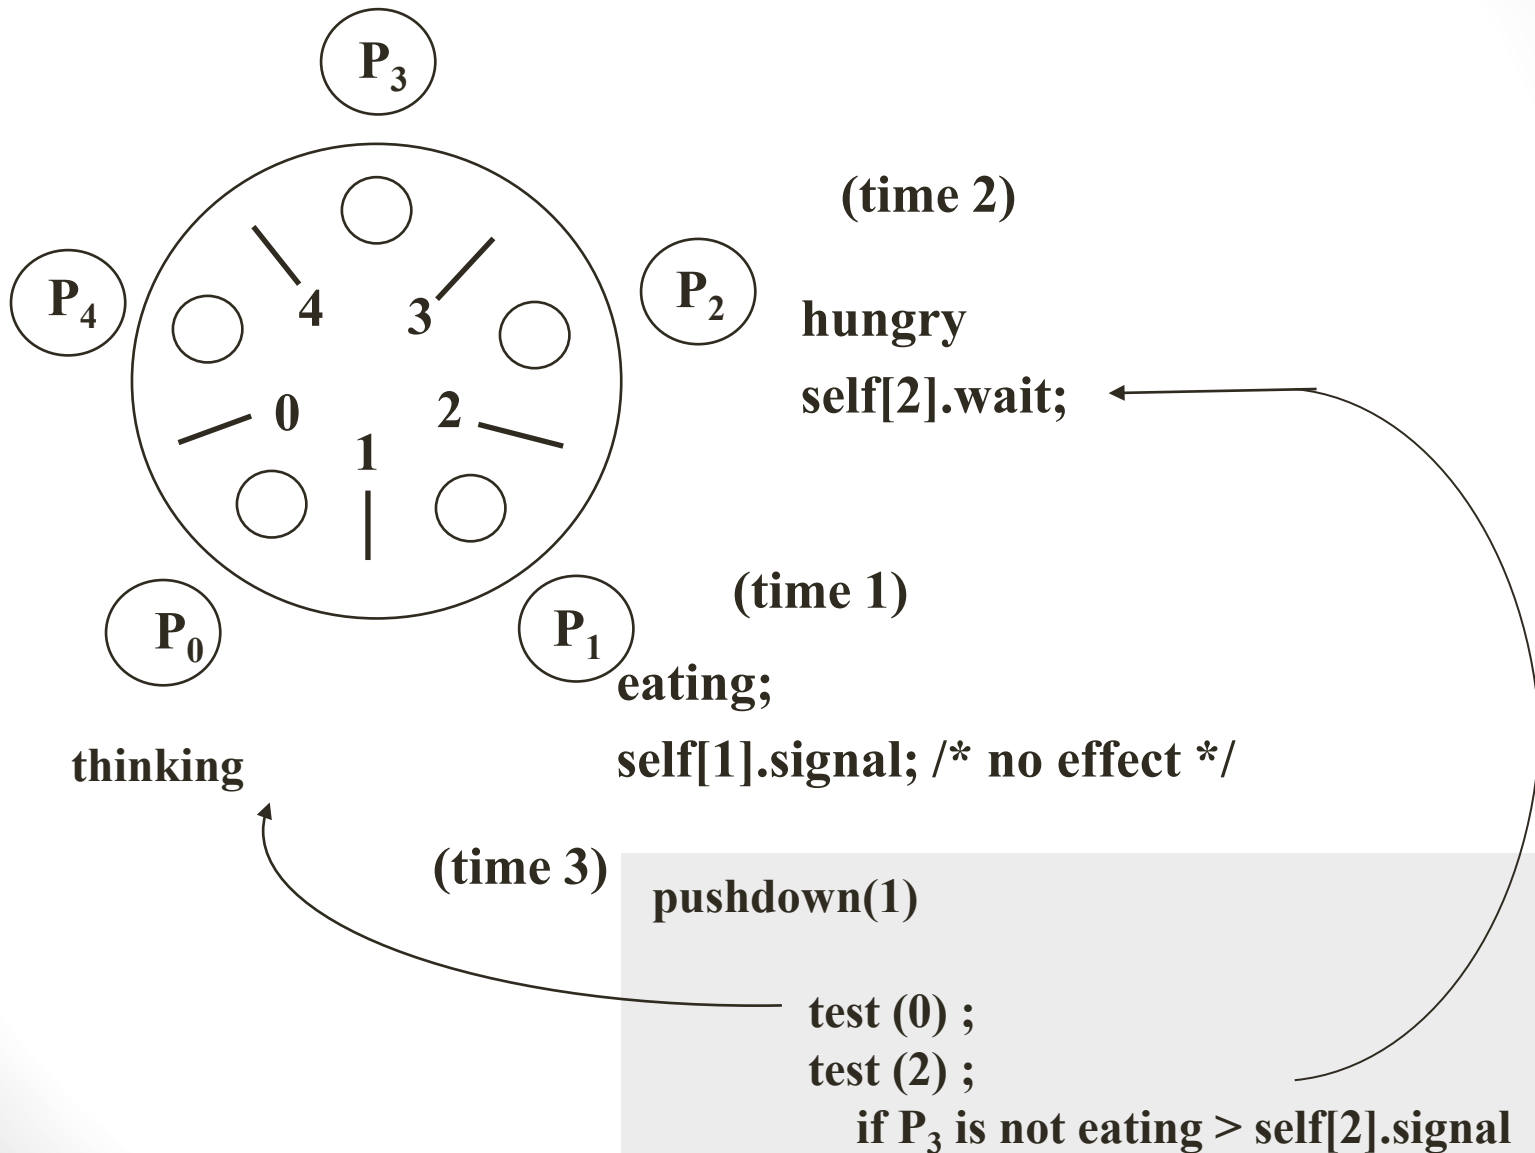
Solution to Dining Philosophers



```
...
void test (int i) { /*測試是否可以拿筷子*/
    if ( (state[(i + 4) % 5] != EATING) && //看旁邊是否正在吃
        (state[i] == HUNGRY) && //自己是否餓了?
        (state[(i + 1) % 5] != EATING) //看旁邊是否正在吃
    ) {
        state[i] = EATING ;
        self[i].signal () ;
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
}
```

An illustration



Monitor Implementation using Semaphores

- A possible implementation of the monitor mechanism using semaphores.
 - A binary semaphore **mutex** (initialized to 1) is provided to ensure mutual exclusion for each monitor.
 - A process must **execute wait(mutex) before entering the monitor** and must execute **signal(mutex)** after leaving the monitor

Monitor Implementation using Semaphores

- We will use the signal-and-wait scheme
 - Since a signaling process must wait until the resumed process either leaves or waits, an additional binary semaphore, **next**, is introduced, initialized to 0.
- The signaling processes can use **next** to suspend themselves. An integer variable **next_count** is also provided to count the number of processes suspended on next.

Monitor Implementation Using Semaphores

- Variables
 - semaphore **mutex** = 1; // 是否可進入monitor
 - semaphore **next** = 0; // 用來強迫P等待 P: signaling processes
 - int next-count = 0; // 紀錄有多少個P在等待next
- For every condition x
 - semaphore x-sem = 0; // 用來強迫Q block
 - int x-count = 0; // 統計有多少個Q被卡住

P either waits until Q leaves the monitor or waits for another condition

Monitor Implementation Using Semaphores

Procedure F

```
wait(mutex);  沒有人在monitor 內活動才可進  
  
...body of F...;  
  
if (next_count > 0)  
    signal(next);  若有其他Proc. 存在，記得先解救他們  
else  
    signal(mutex); 若沒有，則開放monitor讓其他人進來
```

在範例中，**P**正結束工作，接著要**Signal**其他人。
若是有人可以**Singal**，就先把對方叫起來，並且自己等待
若沒有，則開放**monitor**讓其他人進來

Monitor Implementation Using Semaphores – condition x

x.wait

```
x-count++;  Q個數加1
if (next_count > 0) {signal(next)}; 若有P存在，先解救P
else {signal(mutex)} 若沒有，讓其他進來
wait(x_sem);  Q等待condition x 而卡住
x-count--;  Q被救之後，Q數量減1
```

x.signal

```
if (x-count > 0) { 如果有Q被卡住才做
    next_count++; P 個數加1
    signal(x_sem); 解救Q
    wait(next);  P 被迫等待
    next_count--; P被救之後，P數量減1
}
```

在範例中，**P**正結束工作，接著要**Signal**其他人。
若是有人可以**Singal**，就先把對方叫起來，並且自己等待
若沒有，則開放**monitor**讓其他人進來

Monitor

- Process-Resumption Order
 - Queuing mechanisms for a monitor and its condition variables. One simple solution is to use a first-come, first-served (FCFS) ordering ...
- The **conditional-wait** construct can be used to adequate more complicate case
 - `x.wait(c)`
 - **C : priority number**, stored with the name of the process that is suspended.
 - `x.signal()`
 - The process with the smallest priority number is resumed next

A Monitor to Allocate Single Resource

- To illustrate these, consider the **ResourceAllocator** monitor, which controls the allocation of a single resource among competing processes.
 - Each process, when requesting an allocation of this resource, **specifies the maximum time it plans to use the resource.**
 - The monitor allocates the resource to the **process that has the shortest time-allocation request.**

A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
```

```
{
```

```
    boolean busy;
```

```
    condition x;
```

```
    void acquire(int time) {
```

```
        if (busy)
```

```
            x.wait(time);
```

```
            busy = TRUE;
```

```
    }
```

```
    void release() {
```

```
        busy = FALSE;
```

```
        x.signal();
```

```
    }
```

```
    initialization code() {
```

```
        busy = FALSE;
```

```
    }
```

```
}
```

```
R.acquire(t);
```

```
...
```

```
access the resource;
```

```
...
```

```
R.release;
```

time =
the maximum time it
plans to use the resource

**The monitor allocates the
resource to the process that has
the shortest time-allocation
request.**

Note! Defect still possible!

- The monitor concept cannot guarantee that the preceding access sequence will be observed.
 - A process might **access a resource without first gaining access permission** to the resource.
 - A process might **never release a resource** once it has been granted access to the resource.
 - A process might **attempt to release a resource that it never requested**.
 - A process might **request the same resource twice** (without first releasing the resource).

We need to make sure ...

- First, user processes must always make their **calls on the monitor in a correct sequence.**
- Second, we must be sure that **an uncooperative process does not simply ignore the mutual-exclusion** gateway provided by the monitor and try to access the shared resource directly, **without using the access protocols.**
 - These treatment may be possible for a small, static system, it is not reasonable for a large system or a dynamic system. – jmp to Chapter 17. for more detail.

Linux Synchronization

- Linux:
 - Prior to kernel Version 2.6, disables interrupts to implement short critical sections
 - Version 2.6 and later, fully preemptive
- Linux provides:
 - Semaphores
 - atomic integers
 - spinlocks
 - reader-writer versions of both
- On single-cpu system, spinlocks replaced by enabling and disabling kernel preemption

Linux Synchronization

- Atomic variables
`atomic_t` is the type for atomic integer

- Consider the variables
`atomic_t counter;`
`int value;`

<i>Atomic Operation</i>	<i>Effect</i>
<code>atomic_set(&counter,5);</code>	<code>counter = 5</code>
<code>atomic_add(10,&counter);</code>	<code>counter = counter + 10</code>
<code>atomic_sub(4,&counter);</code>	<code>counter = counter - 4</code>
<code>atomic_inc(&counter);</code>	<code>counter = counter + 1</code>
<code>value = atomic_read(&counter);</code>	<code>value = 12</code>

POSIX Synchronization

- POSIX API provides
 - mutex locks
 - semaphores
 - condition variable
- Widely used on UNIX, Linux, and macOS

POSIX Mutex Locks

- Creating and initializing the lock

```
#include <pthread.h>

pthread_mutex_t mutex;

/* create and initialize the mutex lock */
pthread_mutex_init(&mutex, NULL);
```

Acquiring and releasing the lock

```
/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```

POSIX Semaphores

- POSIX provides two versions – **named** and **unnamed**.
- Named semaphores can be used by unrelated processes, unnamed cannot.

POSIX Named Semaphores

- Creating and initializing the semaphore:

```
#include <semaphore.h>
sem_t *sem;

/* Create the semaphore and initialize it to 1 */
sem = sem_open("SEM", O_CREAT, 0666, 1);
```

- Another process can access the semaphore by referring to its name **SEM**.
- Acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(sem);

/* critical section */

/* release the semaphore */
sem_post(sem);
```

POSIX Unnamed Semaphores

- Creating an initializing the semaphore:

```
#include <semaphore.h>
sem_t sem;

/* Create the semaphore and initialize it to 1 */
sem_init(&sem, 0, 1);
```

1.A pointer to the semaphore
2.A flag indicating the level of sharing
3.The semaphore's initial value

Acquiring and releasing the semaphore:

```
/* acquire the semaphore */
sem_wait(&sem);

/* critical section */

/* release the semaphore */
sem_post(&sem);
```

POSIX Condition Variables

- POSIX condition variables are associated with a POSIX mutex lock to provide mutual exclusion:
Creating and initializing the condition variable:

```
pthread_mutex_t mutex;  
pthread_cond_t cond_var;
```

```
pthread_mutex_init(&mutex, NULL);  
pthread_cond_init(&cond_var, NULL);
```

POSIX Condition Variables

- Thread waiting for the condition `a == b` to become true:

```
pthread_mutex_lock(&mutex);  
while (a != b)  
    pthread_cond_wait(&cond_var, &mutex);  
  
pthread_mutex_unlock(&mutex);
```

Once this lock is acquired, the thread can check the condition.

releases the mutex lock, thereby allowing **another thread** to access the shared data and possibly update its value so that the condition clause evaluates to true.

- Thread signaling another thread waiting on the condition variable:

```
pthread_mutex_lock(&mutex);  
a = b;  
pthread_cond_signal(&cond_var);  
pthread_mutex_unlock(&mutex);
```

Note: the call to `pthread_cond_signal()` does not release the mutex lock.

Once the mutex lock is **released**, the signaled thread becomes the owner of the mutex lock and returns control from the call to `pthread_cond_wait()`.

Java Synchronization

- Java provides rich set of synchronization features:
 - Java monitors
 - Reentrant locks
 - Semaphores
 - Condition variables

Java Monitors

- Every Java object has associated with it a single lock.
- If a method is declared as **synchronized**, a calling thread must own the lock for the object.
- If the lock is owned by another thread, the calling thread must wait for the lock until it is released.
- Locks are released when the owning thread exits the **synchronized** method.

Bounded Buffer – Java Synchronization

```
public class BoundedBuffer<E>
{
    private static final int BUFFER_SIZE = 5;

    private int count, in, out;
    private E[] buffer;


    public BoundedBuffer() {
        count = 0;
        in = 0;
        out = 0;
        buffer = (E[]) new Object[BUFFER_SIZE];
    }

    /* Producers call this method */
    public synchronized void insert(E item) {
        /* See Figure 7.11 */
    }

    /* Consumers call this method */
    public synchronized E remove() {
        /* See Figure 7.11 */
    }
}
```


Java Synchronization

- A thread that tries to acquire an unavailable lock is placed in the object's **entry set**:



Java Synchronization

- Similarly, each object also has a **wait set**.
- When a thread calls **wait()**:
 1. It releases the lock for the object
 2. The state of the thread is set to blocked
 3. The thread is placed in the wait set for the object



Java Synchronization

- A thread typically calls `wait()` when it is waiting for a condition to become true.
- How does a thread get notified?
- When a thread calls **`notify()`**:
 - An arbitrary thread T is selected from the wait set
 - T is moved from the wait set to the entry set
 - Set the state of T from blocked to runnable.
- T can now compete for the lock to check if the condition it was waiting for is now true.

Bounded Buffer – Java Synchronization

```
/* Producers call this method */
public synchronized void insert(E item) {
    while (count == BUFFER_SIZE) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }

    buffer[in] = item;
    in = (in + 1) % BUFFER_SIZE;
    count++;

    notify();
}
```

Bounded Buffer – Java Synchronization

```
/* Consumers call this method */
public synchronized E remove() {
    E item;

    while (count == 0) {
        try {
            wait();
        }
        catch (InterruptedException ie) { }
    }

    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    count--;

    notify();

    return item;
}
```

Java Reentrant Locks

- Similar to mutex locks.
- The **finally** clause ensures the lock will be released in case an exception occurs in the **try** block.

```
Lock key = new ReentrantLock();

key.lock();
try {
    /* critical section */
}
finally {
    key.unlock();
}
```

Java Semaphores

- Constructor:

```
Semaphore(int value);
```

- Usage:

```
Semaphore sem = new Semaphore(1);

try {
    sem.acquire();
    /* critical section */
}
catch (InterruptedException ie) { }
finally {
    sem.release();
}
```


Java Condition Variables

- Condition variables are associated with an **ReentrantLock**.
- Creating a condition variable using **newCondition()** method of **ReentrantLock**:

```
Lock key = new ReentrantLock();  
Condition condVar = key.newCondition();
```

- A thread waits by calling the **await()** method, and signals by calling the **signal()** method.

Java Condition Variables

- Example:
- Five threads numbered 0 .. 4
- Shared variable **turn** indicating which thread's turn it is.
- Thread calls **doWork ()** when it wishes to do some work. (But it may only do work if it is their turn.
- If not their turn, wait
- If their turn, do some work for awhile
- When completed, notify the thread whose turn is next.
- Necessary data structures:

```
Lock lock = new ReentrantLock();
Condition[] condVars = new Condition[5];

for (int i = 0; i < 5; i++)
    condVars[i] = lock.newCondition();
```

Java Condition Variables

- Example:
- Five threads numbered 0 .. 4
- Shared variable `turn` indicating which thread's turn it is.

```
Lock lock = new ReentrantLock();  
Condition[] condVars = new Condition[5];  
  
for (int i = 0; i < 5; i++)  
    condVars[i] = lock.newCondition();
```

- Necessary data structures:

```

/* threadNumber is the thread that wishes to do some work */
public void doWork(int threadNumber)
{
    lock.lock();

    try {
        /**
         * If it's not my turn, then wait
         * until I'm signaled.
         */
        if (threadNumber != turn)
            condVars[threadNumber].await();



/**
             * Do some work for awhile ...
             */



        /**
         * Now signal to the next thread.
         */
        turn = (turn + 1) % 5;
        condVars[turn].signal();
    }
    catch (InterruptedException ie) { }
    finally {
        lock.unlock();
    }
}

```

Transactional Memory

- The concept of transactional memory originated in **database** theory.
- A **memory transaction** is a sequence of memory read–write operations that are atomic.
 - As an alternative to traditional locking methods, new features that take advantage of transactional memory **can be added to a programming language**.
- It is only a concept

Transactional Memory

- Consider a function `update()` that must be called atomically. One option is to use mutex locks:

```
void update ()
{
    acquire();

    /* modify shared data */

    release();
}
```

- A **memory transaction** is a sequence of read-write operations to memory that are performed atomically. A transaction can be completed by adding **`atomic{S}`** which ensure statements in **`S`** are executed atomically:

```
void update ()
{
    atomic {
        /* modify shared data */
    }
}
```

Using OpenMP

- OpenMP is a set of compiler directives and API that support parallel programming.

```
void update(int value)
{
    #pragma omp critical
    {
        count += value
    }
}
```

The code contained within the **#pragma omp critical** directive is treated as a critical section and performed atomically.

SUGGESTION! OR OBJECTION?

Let's stop here,

TAKE A BREAK